

CityTech TTP Summer Bootcamp

Data Wrangling

2026-07-07

Agenda

1. Review
2. Follow-Ups
3. Dirty vs Clean Data
4. Wrangling Techniques
5. Practice

Review

What are the most important things you remember from yesterday?

Follow Up: Clean pgAdmin Backups

Restoring a dump onto an existing database fails with "already exists" errors — unless the dump **drops the old objects first**.

In pgAdmin's **Backup** dialog, open the **Query Options** tab and enable:

- **Include DROP DATABASE statement** — adds `DROP` commands so a restore clears the old objects first
- **Include IF EXISTS clause** — makes those drops safe when nothing is there yet (only valid with the option above)
- **Include CREATE DATABASE statement** — if the dump should recreate the database too

Then your `dump.sql` restores cleanly on any machine.

Follow Up: Installing pandas (and Friends)

Libraries must go into the **environment your kernel runs in**.


The safe way from **inside a notebook** installs into the *running* kernel:

```
%pip install pandas
```

Gotcha: a plain `pip install` in the wrong shell hits a **different Python** → `ModuleNotFoundError`.

Follow Up: Two Different Environments

Your terminal and your notebook can be **separate Python environments** — `pip install` lands the package in one, while the kernel looks in the other.



Continue: Squirrel Census EDA

Back to yesterday's notebook: `demos/eda-squirrel/squirrel_eda.ipynb`.

- Reopen it and **Restart Kernel and Run All Cells** for a clean start
- Keep exploring: try `value_counts()` on `age`, `shift`, and `location`
- Notice the rough edges: missing values, a date stored as a number

Those rough edges are our next topic.

Dirty Data vs Clean Data

Real-world data is usually **dirty** — it needs work before you can trust an analysis.

Dirty data has:

- **Missing** values (blanks, NaN)
- **Inconsistent** formats ("NYC" vs "New York")
- **Duplicates**, typos, and wrong **types** (a date stored as a number)
- **Outliers** and impossible values

Clean data is consistent, correctly typed, de-duplicated, and ready to analyze.

Data Wrangling

The broad umbrella term is **data wrangling** (a.k.a. **munging**): turning raw, messy data into a usable form.

It covers a range of steps:

- **Cleaning** — fix missing values, duplicates, wrong types, bad formats
- **Transforming** — reshape, derive new columns, aggregate
- **Combining** — merge and join data from multiple sources

Cleaning is usually the first and biggest part — so we'll start there.

Cleaning: Missing Values

Real data has gaps — blank cells that pandas shows as **NaN**.

- **Find** them: `df.isna().sum()` counts missing per column
- **Drop** rows or columns with `dropna()`
- **Fill** (impute) with `fillna()` — a constant, the mean/median, or forward-fill
- It's a judgment call: dropping loses data, filling makes assumptions

Missing Values: Code

```
# count missing per column
df.isna().sum()

# drop rows missing a fur color
df = df.dropna(subset=["primary_fur_color"])

# or fill missing ages with a placeholder
df["age"] = df["age"].fillna("Unknown")
```

Cleaning: Duplicates

Duplicate rows sneak in from merges, re-exports, or double entries — and they **skew counts and averages**.

- **Find** them: `df.duplicated().sum()`
- **Remove** them: `df.drop_duplicates()`
- Drop by a **subset** of columns when only some define a "duplicate"

Duplicates: Code

```
# how many fully-duplicate rows?  
df.duplicated().sum()  
  
# drop exact duplicate rows  
df = df.drop_duplicates()  
  
# or treat rows as duplicate by a subset of columns  
df = df.drop_duplicates(  
    subset=["hectare", "shift", "date", "hectare_squirrel_number"]  
)
```

Cleaning: Data Types

A column read as the wrong **type** breaks analysis — a date stored as an integer, or numbers stored as text.

- **Check types:** `df.dtypes` (or `df.info()`)
- **Convert** simple casts with `astype()`
- **Parse dates** with `pd.to_datetime()`
- Correct types unlock the right operations (date math, numeric stats)

Data Types: Code

```
# the squirrel "date" is an int like 10142018
df["date"] = pd.to_datetime(df["date"].astype(str), format="%m%d%Y")

# cast a column to a category
df["age"] = df["age"].astype("category")
```

pandas Data Types (dtypes)

dtype	Holds	Like Python's
<code>int64</code>	whole numbers	<code>int</code>
<code>float64</code>	decimals (holds <code>NaN</code>)	<code>float</code>
<code>bool</code>	True / False	<code>bool</code>
<code>string (str)</code>	text	<code>str</code>
<code>category</code>	a fixed set of labels	no direct one

category has no Python counterpart: like an **enum**, it stores each label once and cells as codes — ideal for few-valued columns (fur color, shift): less memory, faster grouping.

Cleaning: Standardize Formats

The same thing written many ways ("NYC", "nyc", " New York ") counts as **different values** unless you normalize it.

- **Trim** whitespace and fix **casing**: `.str.strip()` , `.str.lower()`
- **Unify** labels and codes with `replace()` or a mapping
- Consistent formats make grouping and joining reliable

Standardize: Code

```
# trim whitespace and normalize casing
df["primary_fur_color"] = df["primary_fur_color"].str.strip().str.title()

# map inconsistent labels to one form
df["shift"] = df["shift"].replace({"am": "AM", "pm": "PM"})
```

Cleaning: Outliers

Outliers are extreme or impossible values — a negative age, a \$1M coffee. They can distort means and models.

- **Spot** them with `describe()`, a histogram, or the min/max
- Decide: **remove**, **cap** (clip to a limit), or **flag** and keep
- Not every outlier is an error — sometimes it's the interesting part

Outliers: Code

```
# spot the range
df["hectare_squirrel_number"].describe()

# keep only plausible values
df = df[df["hectare_squirrel_number"] <= 20]

# or cap extremes instead of dropping
df["hectare_squirrel_number"] = df["hectare_squirrel_number"].clip(upper=20)
```

Cleaning: Rename Columns

Clear, consistent column names make everything downstream easier to read and type.

- **Rename** specific columns with a mapping: `rename(columns=...)`
- Standardize a **convention** (e.g. lowercase `snake_case`)
- Do it early so the rest of your code reads cleanly

Rename: Code

```
# rename specific columns
df = df.rename(columns={
    "primary_fur_color": "fur_color",
    "hectare_squirrel_number": "squirrel_no",
})

# or lowercase every column name
df.columns = df.columns.str.lower()
```

Code Along: Mural Census

Your turn — start wrangling a messy dataset.

- Data: `demos/eda-murals/` — `murals.csv` + `condition_reports.csv` (5,000 rows each)
- Start your own `practice.ipynb` in that folder
- Clean it up: **rename, fix types, missing, duplicates, standardize, outliers**
- We'll **bin** and **merge** tomorrow

(Fake data, generated for practice — messy on purpose.)

Today's Exit Ticket

- [] Explain the difference between **dirty** and **clean** data
- [] Explain **data wrangling** (and how **cleaning** fits under it)
- [] Clean a DataFrame: **missing values, duplicates, types, formats, outliers, renaming**