

CityTech TTP Summer Bootcamp

Describing Data and Telling Stories

2026-07-08

Agenda

1. Review
2. Follow-Ups: Standard Deviation and Quartiles
3. Writing Pythonic Code
4. Transforming and Combining Data
5. The Garden of Forking Paths
6. Practice

Review

What are the most important things you remember from yesterday?

Follow Up: Why Standard Deviation?

The **mean** alone can hide a lot - two datasets can share a mean yet look nothing alike:

- A: [68, 70, 71, 71] - mean **70**, tightly clustered
- B: [40, 60, 80, 100] - mean **70**, widely spread

We need one number for **how spread out** the data is. That's **standard deviation**, written σ (the Greek letter *sigma*).

Standard Deviation: The Math

$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}$$


1. Find the mean μ
2. Take each value's **distance from the mean**, and **square** it
3. **Average** the squared distances → the **variance**
4. **Square root** the variance → the standard deviation σ

Standard Deviation: In Plain Words

Standard deviation is the **typical distance** of a value from the mean.

- **Small σ** → values huddle **close** to the mean (consistent)
- **Large σ** → values are **spread far** from the mean (variable)
- It's in the **same units** as the data (dollars, inches, points)
- Rule of thumb (bell-shaped data): about **68%** of values fall within **one σ** of the mean


Standard Deviation: Picture It



Same mean, different spread - the **narrow** curve has a **small σ** , the **wide** curve a **large σ** .

The 68-95-99.7 Rule

For bell-shaped data, σ predicts how much falls within each band of the mean:



Follow Up: Why Quartiles?

The mean and σ get **distorted by outliers** and skew. Quartiles don't.

- Salaries (thousands): `[30, 32, 35, 38, 40, 500]`
- The **mean** (~112) is dragged up by one value; the **median** (~36) isn't
- Quartiles describe spread by **position**, so they're **robust** to extremes

Quartiles: The Definition

Sort the data, then split it into **four equal parts** at three cut points:

- **Q1: the 25th percentile** (a quarter of the data is below it)
- **Q2: the 50th percentile, the median**
- **Q3: the 75th percentile** (three quarters below it)

IQR (interquartile range) = **Q3 – Q1** - the spread of the middle **50%**.


Quartiles: In Plain Words

Quartiles mark where a value stands relative to everyone else.

- **Q1**: the bottom **25%** are below it
- **Q2 / median**: right in the **middle**
- **Q3**: only the top **25%** are above it
- The **IQR** is the range of the "typical" middle half – narrow means clustered, wide means spread

Quartiles: Picture It


The **box plot** (bottom) summarizes the raw **data** (top) - for a symmetric and a skewed dataset:



Box = **Q1-Q3**, line = **median**, whiskers reach the rest, dots beyond = **outliers**.

Same Mean and σ , Different Skew

Both distributions have the **same mean and same σ** - one symmetric, one skewed - so those numbers can't tell them apart. The **box plots** expose the skew:



One Obvious Way

Python favors **one clear, obvious way** to do a common task - not a dozen clever ones.

- The aim is code that's **readable** and instantly recognizable to any Python developer
- Write with the language's grain: lean on **built-in features** instead of reinventing them
- Code that does this is called **Pythonic**

The Zen of Python

Python's guiding principles ([PEP 20](#)) - try `import this`:

```
Beautiful is better than ugly.  
Explicit is better than implicit.  
Simple is better than complex.  
Flat is better than nested.  
Readability counts.  
There should be one-- and preferably only one --obvious way to do it.
```

JS vs Python: Iterating

JavaScript

```
items.forEach((item, i) => {  
  console.log(i, item);  
});
```

Python

```
for i, item in enumerate(items):  
    print(i, item)
```

JS vs Python: Build a List

JavaScript

```
const nums = [-2, 3, -1, 4];  
const result = [];  
for (const x of nums) {  
  if (x > 0) result.push(x * x);  
}
```

Python

```
nums = [-2, 3, -1, 4]  
result = [x * x for x in nums if x > 0]
```

JS vs Python: Truthiness

JavaScript

```
if (items.length > 0) { ... }  
if (items.length === 0) { ... }
```

Python

```
if items:  
    ...  
if not items:  
    ...
```

JS vs Python: Two Lists at Once

JavaScript

```
for (let i = 0; i < names.length; i++) {  
  console.log(names[i], scores[i]);  
}
```

Python

```
for name, score in zip(names, scores):  
  print(name, score)
```

Transforming: Binning

Binning groups a continuous number into labeled **buckets** - ages into "child / adult / senior", or counts into "low / high".

Why bin?

- **Simpler to read:** a few groups instead of hundreds of distinct values
- **Reveals patterns:** trends per bucket stand out more than raw numbers
- **Enables grouping:** now you can `value_counts()`, group, and chart by category
- **Human-meaningful:** "senior" says more than "age 71"

Trade-off: binning **loses precision** - bin when the groups matter more than exact numbers.

Binning: Code

```
# bucket squirrel counts into low / medium / high
df["density"] = pd.cut(
    df["hectare_squirrel_number"],
    bins=[0, 3, 8, 100],
    labels=["low", "medium", "high"],
)

df["density"].value_counts()
```

Combining: Merge / Join

Merging combines two DataFrames by matching rows on a shared **key** column - just like a SQL join.

- `df.merge(other, on="key")` lines up rows where the key matches
- `how=` controls what to keep: **inner** (matches only), **left**, **right**, **outer**
- Great for **enriching** data - pull in details from another table

Merge: Code

```
# a lookup table: hectare -> zone
hectares = pd.DataFrame({
    "hectare": ["01A", "01B", "02A"],
    "zone":    ["north", "north", "south"],
})

# attach each squirrel's zone by matching on `hectare`
df = df.merge(hectares, on="hectare", how="left")
```

Wrangling Is Full of Choices

Every wrangling step is a **judgment call** - reasonable analysts would choose differently:

- Drop rows with missing data, or fill them (and with what - mean? median? zero?)
- Which outliers are **errors** to remove vs. **real** values to keep?
- Where do the bin edges go? "Adult" starts at what age?
- Which rows and columns to keep or filter out?

None of these are purely objective - each one **shapes the result**.

The Garden of Forking Paths

Statistician **Andrew Gelman**'s idea: even an honest analyst walks a **branching maze** of reasonable choices.

- Each decision is a **fork**; different forks lead to different results
- You might run only **one** analysis - but it was one path out of many
- So a striking finding can be an **artifact of the path** you took, not a real effect

(The name is from a Borges story - a labyrinth of forking paths.)

Same Data, Many Stories

- **Data is (relatively) objective:** the recorded facts
- **Data analysis tells a story:** shaped by the choices you make
- **Data analysis can tell many stories:** different honest paths, different conclusions
- And not every path is honest - data can be **intentionally misrepresented** to sell a conclusion

Be **transparent** about your choices, and **skeptical** of any single story - including your own.

Continuing: Mural Census Code Along

Pick up your `practice.ipynb` from yesterday and finish the job.

- Data: `demos/eda-murals/` (you started cleaning it yesterday)
- Now **bin** a column (e.g., mural size) into categories
- Then **merge** `murals` + `condition_reports` on `m_id`
- Explore the combined data with `value_counts()` and `groupby()`

(Fake data, generated for practice - messy on purpose.)

Pair Programming: Fridge Census

Pair up and wrangle a fresh messy dataset together.

- Data: `demos/eda-fridge/` - `fridges.csv` + `stocking_log.csv` (5,000 rows each)
- One **driver** (types), one **navigator** (guides) - then switch
- Apply every technique: **rename, fix types, missing, duplicates, standardize, outliers, bin**
- Then **merge** the two tables on `f_id`

(Fake data, generated for practice - messy on purpose.)

Pair Programming: NYC Open Data

Pair up and take a **real** dataset from raw to insight – start to finish.

- **Find:** browse [NYC Open Data](#) and pick a dataset that interests you
- **Explore:** `head()`, `info()`, `describe()`, `value_counts()` – get a feel for it
- **Clean:** fix types, missing values, duplicates, outliers, inconsistent formats
- **Insight:** ask a question and try to answer it

Swap **driver** and **navigator** as you go, and be ready to share one finding.

Today's Exit Ticket

- [] Explain what **standard deviation** and **quartiles** each tell you about spread
- [] Explain why two datasets can share a **mean and σ** but differ in **skew**
- [] Explain what makes code **Pythonic**, and give an example
- [] Use **binning** and **merge / join** to transform and combine data
- [] Explain the **garden of forking paths** - why analysis choices shape the story
- [] Examine a **real-world dataset**: explore it, clean it, and draw an insight